

Classic Synchronization Problems and Monitors

Operating Systems and Multiprogramming

Lecturer: Prof. Eric Jul

29-Feb-2016



Present by: **Marjan Mansourvar**

Contents



Short recap on last lecture

- ☐ The Critical-Section Problem
- ☐ Solutions to the critical-section problems
- ☐ Requirements



Software solutions to the critical section problem

- ☐ Peterson's Solution
- ☐ Semaphores



Classic Synchronization Problems

- ☐ The Dining- Philosophers problems
- ☐ The Readers- Writers Problem
- ☐ The bounded-Buffer Problem

Contents

◆ **Monitors**

- ☐ Concept and Invention
- ☐ How to: Signal
- ☐ Monitoring the Philosophers
- ☐ Monitors in Java

◆ **Synchronization Examples**

- ☐ Transactional Memory

◆ **THE multiprogramming system**

The Critical-Section Problem

❖ Critical Section:

- ❑ A system consisting of n processes (p_0, p_1, \dots, p_{n-1}), each process has a segment code called critical section

❖ Critical section problem

- ❑ is to design a protocol that processes can use to cooperate

General structure of a process P_i

Do{
Entry section

Critical section

Exit section
Remainder section
}

The Critical-Section Problem

❖ 3 Requirements :

❑ **Mutual exclusion** (Protection of critical sections)

- ❖ If process P_i is executing in its critical section then no other process can be executing in their critical section .

❖ **Progress**

- ❑ The decision as to who can enter the critical section must not be postponed indefinitely.
- ❑ (when no process is in its critical section and several processes are requesting to enter their critical section, who should decide who will enter the critical section?)

❖ **Bounded waiting**

- ❑ Exists a bound or limit on the number of times that the process are allowed to enter their CSs

Peterson's Solutions

- ❖ Restricted to 2 processes
- ❖ Requires 2 processes to share 2 data items:

int turn;

Boolean flag [2];

Turn==i; means **Pi** is **allowed to execute** in its CS

Flag[i]==true; means **pi** is **ready to enter** its CS

- ❖ Mutual Exclusion is preserved (Since the value of turn can be either 0 or 1)
- ❖ The progress requirement is satisfied
- ❖ The bounded-waiting is met

```

P0: do {
    flag[0] = true;
    turn = 1;
    while (flag[1] && turn == 1) do skip;

```

```

    // critical section

```

```

    ...

```

```

    // end of critical section

```

```

    flag[0] = false;

```

```

    Reminder section

```

```

P1:    do {
    flag[1] = true;
    turn = 0;
    while (flag[0] && turn == 0) do
        skip;
    {

```

```

        // critical section

```

```

        ...

```

```

        // end of critical section

```

```

        flag[1] = false;

```

```

        Reminder section

```

Semaphores (Dijkstra, 1965)

- ❖ A high-level method for processes to synchronize activities
- ❖ A semaphore is an integer variable
- ❖ is accessed only through 2 standard operation :
wait() and **signal ()** (P or V)

```
Wait (s) {  
  while (s <= 0)  
    ; // busy  
  s--;  
}
```

```
Signal (s) {  
  S++;  
}
```




Going to the second parts

Classic Problems in Synchronization

- ❖ **The Dining Philosophers**
- ❖ **The Readers-Writers Problems**
- ❖ **The Bounded-Buffer Problem**


The Dining –Philosophers Problem

❑ It was originally formulated in 1965 by Edsger Dijkstra

A simple example showing the need to allocate several resources among several processes without deadlocks or starvation.

Problem Statement

- ❖ Five philosophers, sitting around a table, sharing five forks.
- ❖ The philosophers alternates between thinking, becoming hungry and eating.
- ❖ When hungry, they try to pick up a fork on his left and right.
- ❖ Then he goes back to solving the problems of humanity.



Feeding the Philosophers with Semaphores


Represent each fork with a semaphore

- ❑ Try to grab a fork by execution a wait() operation on the semaphore
- ❑ Release her fork by exciting the signal ()

```
Semaphore fork [5];  
do {  
    wait (fork [i]);  
    wait (fork [(i+1) %5]);  
    .....  
    / eat for awhile/  
    .....  
    signal (fork[i]);  
    signal (fork [(i+1) %5 ]);  
    .....  
    / think for awhile/  
} while (true);
```

Feeding the Philosophers with Semaphores


- ❖ Does this work?
- ❖ No. Risk of **deadlock**
- ❖ The Starving Philosophers!



Solutions to the Starving Philosophers

- ❖ Allow at most four philosophers to be sitting simultaneously at the table .
- ❖ Allow a philosopher to pick up his fork only if both forks are available.
- ❖ Use an asymmetric solution(odd-number picks up first left and then right fork, while even-numbered picks up right fork and then left fork)

The Readers-Writers Problems



- ❑ A database is to be shared among several concurrent processes.
- ❑ Some of those want only to read (readers are allowed to read concurrently.)
- ❑ Some of them want to update. (writers must acquire exclusive access.)
- ❑ Has several Variations , all involving priority

The Simple Semaphore Solution

- ❑ Count the number concurrent readers (counter protected by binary semaphore).
- ❑ First reader blocks future writers.
- ❑ Last reader unblocks writers.
- ❑ Ensure mutual exclusion amongst writers using a **write** semaphore.

Readers process share data structure

Initialization

```
Semaphore rw_mutex = 1;
```

```
Semaphore mutex = 1;
```

```
int read_count = 0;
```


The Readers-Writers Problems

Reader

```

    Do {
        Wait (mutex);
        Read_count++;
        If (read_count == 1)
            Wait(rw_mutex);
        Signal(mutex);
        .....
        /*reading is performed*/
        .....
        Wait(mutex);
        Read_count--;
        If(read_count== 0)
            Signal(rw_mutex);
        Signal(mutex);
    } while (true);

```

Initialization

```

Semaphore rw_mutex =1;
Semaphore mutex = 1;
int read_count = 0;

```

Writer


```

    do {
        Wait (rw_mutex);
        .....
        /* writing is performed */
        .....
        Signal (rw_mutex)
    } while ( true) ;

```

The Bounded-Buffer Problem

- ❑ The producer–consumer problem also known as the bounded-buffer problem
- ❑ Suppose a circular buffer with two pointers ***in*** and ***out*** to indicate the next available position for depositing data and the position that contains the next data to be retrieved.
- ❑ Bounded buffer: size 'N'
(Access entry 0... N-1, then “wrap around” to 0 again)
- ❑ Producer process writes data to buffer
(Must not write more than 'N' items more than consumer “ate”)
- ❑ Consumer process reads data from buffer
(Should not try to consume if there is no data)



Solving with semaphores

```
int n
```

```
semaphore mutex = 1; /* for mutual exclusion*/
semaphore empty = N; /* number empty buffer entries */
semaphore full = 0; /* number full buffer entries */
```

Producer

```
do {
    .....
    /*produce an item in next produced*/
    wait(empty);
    wait(mutex);

    .....
    /*add next produced to the buffer*/

    .....
    signal(mutex);
    signal(full);
} while (true);
```

Consumer

```
do {
    wait(full);
    wait(mutex);

    .....
    //remove an item from buffer*/
    .....
    signal(mutex);
    signal(empty);

    .....
    /

    * consume the item in next-
    consumed*/

    .....
} while (true);
```

What's wrong with Semaphores?

- ❖ They are essentially shared global variables.
- ❖ Access to semaphores can come from anywhere in a program.
- ❖ They serve multiple purposes (mutual exclusion , scheduling constraint, ...)
- ❖ There is no control or guarantee of proper usage.

❖ **Solution: use a higher level primitive called monitors**

Hoare Monitors

A synchronization construct that allows threads to have both mutual exclusion and the ability to wait (block) on certain conditions.

Invented by C. A. R. Hoare (aside: read up on CSP; it's awesome) in 1974. First implemented by the late Danish-American Per Brinch Hansen in 1993 as part of his Concurrent Pascal language.

Basically a mutex and a number of condition variables.

In terms of object-oriented programming (where we most often see monitors), it is a thread-safe class, object or module that uses implicit mutual exclusion to allow access to a method or variable by concurrent threads.

Hoare Monitors

- ❑ Only one thread can execute any monitor procedure at any time (the thread is “in the monitor”)
 - If a second thread invokes a monitor procedure when a first thread is already executing one, it blocks
- ❑ The only operations can be invoked on a condition are wait () and signal ().

❑ **X.wait ();**

- The process invoking it is suspended until another process invoke

❑ **X. signal () ;**

- This process resume exactly one suspended process.

Monitors: Key Features

Higher level construct than semaphores.

A **package** of grouped procedures, variables and data

Processes can only call procedures within a monitor

Can be built into **programming languages**.

Synchronization is enforced by the compiler.

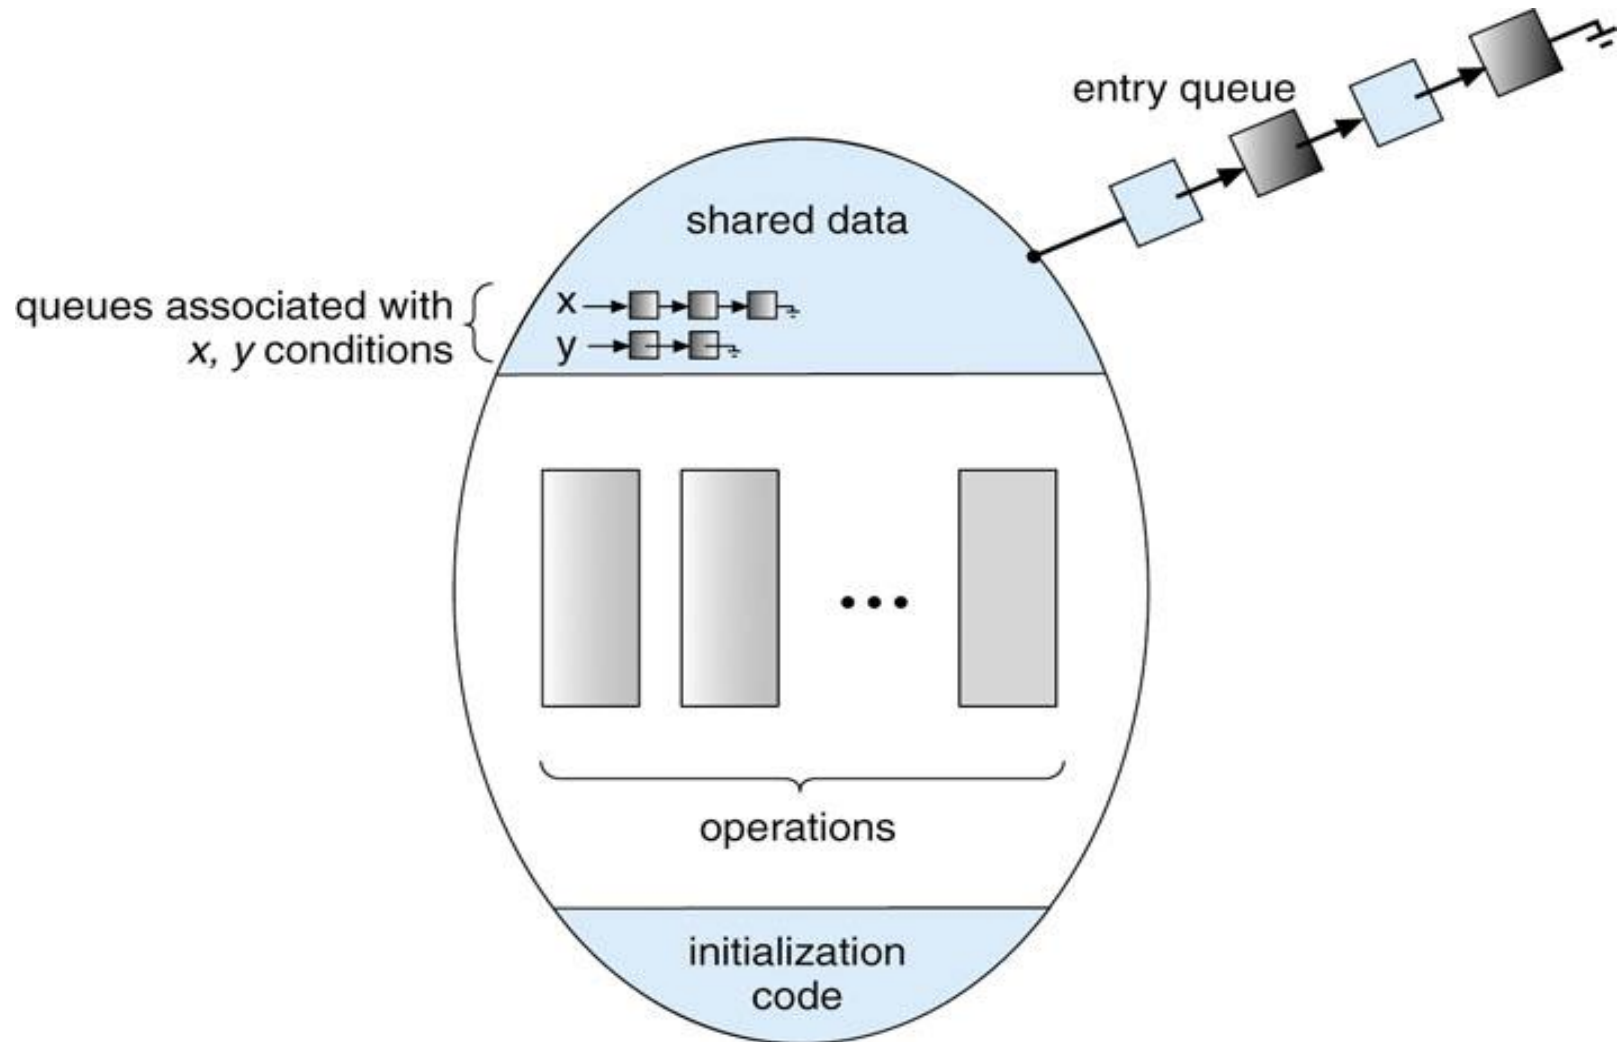
Only **one process** allowed inside the monitor at any
time.

Wait and signal on condition variables.



Perfect

View of a Monitor



Are you ready to continue ????



Monitoring the Dining-Philisophers

- ❖ Monitor concepts by presenting a deadlock-free solution
- ❖ Imposes the restriction that a philosopher may pick up the forks only if both of them are available.
- ❖ Data structure :
- ❖ `Enum{THINKING, HUNGRY, EATING} state [5]`
- ❖ Philosopher i can set the variable `state[i] = EATING` only if 2 neighbors are not eating: `(state[(i+4)%5] != EATING)` and `(state[(i+1)%5] != EATING)`.

Monitoring the Dining-Philosophers

```
monitor DiningPhilosophers {
enum {THINKING , HUNGRY , EATING}
    state[5];
    condition self[5];
```

```
    void pickup(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }
```

```
    void putdown(int i) {
        state[i] = THINKING;
        test(LEFT(i));
        17 test(RIGHT(i));
    }
```

```
    void test(int i) {
        if ((state[LEFT(i)] != EATING) &&
            (state[i] == HUNGRY) &&
            (state[RIGHT(i)] != EATING)) {
            state[i] = EATING;
        }
```

```
Initialization_code ()
{ for (int i = 0; i < 5; i++)
    state[i] = THINKING; }
LEFT(i) = ((i+1) % 5)
define RIGHT(i)= ((i+4) % 5)
```

Resuming Processes within a Monitor

If several processes are suspended on condition X, and X.signal() operation is executed, how we determine which process should be resumed?

- ❑ Use a first-come, first-served (FCFS) ordering
(The process has been waiting the longest is resumed first .
- ❑ **Conditional-wait** construct can be used.

x.wait(c);

c is an integer expression called a **priority number**

Monitors in Java

An example of a language that includes monitor support directly:

- ❑ Each object has an associated lock.
- ❑ Declare methods synchronized to enable mutual exclusion.
- ❑ Condition variables replaced by **wait()** and **notify()** methods.
- ❑ Java API supports semaphores, condition variable, and mutex lock in the **Java.util.concurrent** package .

```
Public class SimpleClass {  
    .....  
    Public synchronized void  
    safeMethod () {  
        ....  
        /* Implamnetion of safeMethod() */  
        .....  
    }  
}
```

Transactional Memory

The main concept of **transaction Memory** originated in database theory

Provide a strategy for process synchronization .

A **transaction Memory** is a sequence of memory read-write operations that are atomic.


If all operations in a transaction are completed the memory transaction is committed.

Otherwise the operations must be rolled-back.

The **benefits of** transaction Memory can be obtained through features added to a programming language.

An example :

```
Void update()
{
    Acquire ();
    /* modify share data*/
    Release();
}
```



We add the construct `atomic{s}` (ensures that the operations in S execute as a transaction.

Advantages:

- The transactional memory system- not the developer- is responsible for guaranteeing
 - deadlock is not possible
 - Identify concurrent read access to the share variable

```
Void update ()
{
    atomic {
        /* modify share data*/
    }
}
```



Be PatientPeak is near

THE multiprogramming system

THE multiprogramming system was a computer operating system designed by a team led by **Edsger W. Dijkstra**, described in monographs published in 1968.

The THE system was primarily a **batch system** that supported multitasking;

The THE system apparently introduced the first forms of software-based **memory segmentation**

freeing programmers from being forced to use actual physical locations on the drum memory.

THE multiprogramming system

using a modified **ALGOL** compiler (the only **programming language** supported by Dijkstra's system) to "automatically generate calls to system routines.

The code of the system was written in **assembly language** for the Dutch Electrologica X8 computer.

This system, THE, sported many important features found in modern operating systems such as multiprocessing, virtual memory, resource sharing, and standardized buffered I/O.

The design of the THE multiprogramming system is significant for its use of a layered structure, in which "higher" layers only depend on "lower" layers:

Design

Layer 0 was responsible for the multiprogramming aspects of the operating system. It decided which process was allocated to the CPU, and accounted for processes that were blocked on semaphores.

Layer 1 was concerned with allocating memory to processes.

Layer 2 dealt with communication between the operating system and the console.

Layer 3 managed all I/O between the devices attached to the computer. This included buffering information from the various devices.

Layer 4 consisted of user programs.

Layer 5 was the user (as Dijkstra notes, "not implemented by us").

Summary

- ❑ A look on classical synchronization problems
 - The Dining- Philosophers problems
 - The Readers- Writers Problem
 - The bounded-Buffer Problem
- ❑ The monitor concepts
- ❑ THE multiprogramming system



Thank You !

Computer science

That's all for today

